

COMP-421 Database Systems, Winter 2020

Project 3: Writing your Application

Due Date March 25, 11:59pm

This is the last deliverable for your database application project. If you end up adding more data for this project deliverable, that is ok, as long as you do not EXCEED the record limits imposed in project 2.

1. (15 Pts) Write one stored procedure to perform operations on your project database. It should be nontrivial, illustrating a feature or features such as local variables, multiple SQL statements, loops etc. It should also involve a cursor. The stored procedure should use one or more parameters in a significant way. We encourage you to be imaginative. However, here are some sorts of things you might try if you can't think of something more interesting:
 - Compute some aggregate value from a relation and use that value to modify values in that or another relation.
 - Create a new relation and load it with values computed from one or more existing relations.
 - Enforce a constraint by searching your database for violations and fixing them in some way.

Hand in a listing of your programs and scripts showing them working. You should demonstrate that the programs had their intended effect by querying (before and after) some relation of your project database that was changed by the program. These queries may be included in the file that holds your programs for convenience.

2. (50 Pts) Write a user-friendly application program for your database in Java (You may use another language like Python as long as you can sort out the connectivity to the database). There is no need for a fancy interface. For example, a menu printed via simple console I/O is ok.

Your program should consist of a loop in which:

- A list of at least five alternative options is offered to the user. An additional alternative should be quit.
- The user selects an alternative.
- The system prompts the user for appropriate input values.
- The system accesses the database to perform the appropriate queries and/or modifications.
- Data or an appropriate acknowledgment is returned to the user.

Your program should follow the following guidelines.

- Your options should include both queries and modifications.
- Some of your options should contain more than one SQL statement.
- Your program must handle errors appropriately. For Java, catch exceptions and print the error messages. Ensure that your program terminates gracefully (after closing any connections) even in the case of a database/SQL exception. If we notice your programs are leaving open hundreds of connections that are piling up at the database end and blocking resources, you might get penalized.

For example, if your project were about skaters and competitions.

- Look up whether a skater participates in a certain competition by skater name.
- Enroll a skater S in a competition C. If the rating level is below 3, S cannot enroll in any competition. If it is between 3 and 6, S can enroll in regional competitions only, if its is between 7 and 9, he/she can enroll in regional and national levels, and only with a skating level of 10 can S enroll in all types of competitions. If S is not qualified for the competition C, return a list of alternative competitions for which the S has the minimum rating level and which are close to C in terms of the date.

- A competition is cancelled: find all skaters participating and replace the participation with a competition close in time to the cancelled competition.
- Add a new skater.
- Increase the rating of skaters that were among the first 5 in at least 2 competitions of the highest level they can participate.
- ...
- Quit

Hand in your program and a script showing the program running. Each of the options should be exercised at least once in your script.

Make sure to include screen shots of the program executing various options(similar to Project 2 deliverable). **You may skip the screen shots in your project report for this question if you had instead decided to do a demo to the TA.**

3. (10 pts.) In class we discuss indexes that help to speed up queries. You can create and drop an index using:

```
CREATE INDEX <IndexName> ON <RelName>(<Attribute List>);
e.g., CREATE INDEX skatersname ON Skaters(sname);
DROP INDEX <IndexName>;
```

Statements for more sophisticated indexes (unique, clustered etc.) can be found in the lecture notes and also in the DBS manuals.

Create at least two useful indexes for your project database. Do not build indexes on primary keys and unique constraints. Database systems usually create indexes on these attributes automatically as they need them to check the uniqueness property. For each of the created indexes indicate why this index is useful by describing which application relevant queries would execute quicker

4. (15 Pts.) In the first lecture we mentioned about the role data visualization plays in analyzing data and decision making. Although we will not cover this in the course, nevertheless we will take a small peek into this aspect. For this purpose, you will produce two charts that visualizes some important aspect of your application from the data. Some examples are:

- (a) Overall sales per day/month,...etc.
- (b) Number of user posts per day,...etc.
- (c) Average money spent by a user per month for the top 10 spending users in your store.

Make sure that the two charts do not portray similar data. For example, if one of the charts is sales of a particular product per day, the other should not be sales of another product per day or sales of a product per month, etc...

Choose a chart so that it does not create a “clutter”. Remember, if you cannot read it, there is no point in visualizing it. For example, if you have sales data for an entire year, and decide to produce a chart for sales of each day, you will end up with a clutter on the x axis, and will be unable to read anything. Instead you could chose to just show sales for the last one week, or sales per month.

For this work, you can export your data into excel/google spread sheets.

Almost any desktop client you use to run SQL queries against the database (such as pgadmin, SquirrelSQL, etc..) can export data from a SQL query into CSV format.

If you are using database clients installed in the comp421 server, you have the following options also.

In Postgres, you can export data from a SQL query into CSV format (which can then be loaded to excel/google spread sheet) using the following command at the psql prompt.

```
\COPY (SELECT * FROM skaters) TO skaters.csv WITH CSV
```

For db2 prompt, you can do the same as follows.

```
EXPORT TO result.csv OF DEL MODIFIED BY NOCHARDEL SELECT * FROM skaters
```

You will find the links to two short tutorials in mycourses (under assignment information folder) on how to create pivot charts in Excel and Google spread sheets.

For each chart, turn in the (i) SQL used to generate the data, (ii) a JPG/PNG image of the chart (readable resolution) (iii) The Excel / Google spreadsheet you did the work on.

5. (10 pts.) For creativity points, you may explore any one of these (some are topics not covered in class).

- Triggers - create a trigger that does something interesting for your application logic - give a brief description of its purpose.
- A sophisticated GUI (instead of the console based UI for question 2, but does the same operations)
- An extra stored procedure (it should do something totally different from the one required in this deliverable)
- Do the visualization on your own application (this can be a different standalone application independent of your main application). For example you could use Python matplotlib (you may have to load data using CSV to your program) or javafx (should be able to connect using JDBC to database and read data directly without having to unload it to CSV).

If something else interests you for creativity points, come talk to me.

If you are doing a GUI for creativity points, it is highly recommended to sign up for a demo slot. As stated before, if you are doing a demo for question 2, then you do not have to include the screen shots of the program execution for question 2. You still need to turn in your programs/scripts.

You should still turn in your project report by the due date.

Note : The project design was taken and adjusted from a project description of the CS145 Stanford database course.